

# Divide & Conquer Street-Routing — Method & Experiments 1–3

Version: Oct 2025

Codebase anchors: `divideconquer.py`, `experiment1.py`, `experiment2.py`, `experiment3.py`

---

## 1) Executive summary

- **Goal.** Build practical city routes that follow *real streets*, starting at a depot and ending at another, while covering many delivery points.
  - **Core idea (Divide & Conquer).**
    1. **Divide** the delivery set with a **quad-tree** into small, spatially compact regions.
    2. **Conquer** each region via a **greedy nearest-neighbor** mini-tour on the street network (no straight-line cheats).
    3. **Combine** regions into one city-wide route by **stitching shortest street paths** between region entries and exits.
  - **Experiment 1 — Clean S→N path.** Corridor-aware; selects only regions near the shortest S→N street path; visits regions south→north.
  - **Experiment 2 — Multi-depot.** Repeats the clean-path method across several start→end depot pairs and compares runtime/coverage/memory.
  - **Experiment 3 — Priority-only (no corridor).** Picks **top-scoring regions** by delivery **priority/size**, independent of corridor distance; still orders them south→north to keep a clean narrative.
- 

## 2) Method (what the router actually does)

### 2.1 Inputs

- **Street graph  $G$**  (OSMnx/NetworkX), nodes have lon/lat; edges carry length (meters).
- **Delivery points** CSV with columns: `lat`, `lon`, `priority` (1=high, 2=medium, 3=low). They are snapped to nearest graph node.
- **Depots** (`S`, `N`, `E`, `W`, `C`) with fixed lat/lon.

### 2.2 Divide — Quad-tree

- Recursively split the bounding box of deliveries into 4 quadrants until each region has  $\leq M$  points (e.g., 15–20) or `max_depth` is reached.
- For each region store: points subset, centroid, bounds rectangle, depth, size.

## 2.3 Conquer — Regional mini-tour (street-aware)

- For a chosen region, pick an **entry node** (closest by **graph distance**) to the current position using a single-source shortest path (Dijkstra) from the current node.
- Inside the region, run **greedy nearest-neighbor** on **graph distances** to visit all deliveries. Expand the ordered nodes into a **street polyline** using `nx.shortest_path` between consecutive nodes.

## 2.4 Combine — Stitch regions into one route

- Connect: previous region's last node → next region's entry node via a **street shortest path**; repeat, then end at the terminal depot.

**Guarantee:** Every hop is a path embedded in the street graph, so the final polyline is *non-jumpy* and map-clean. (If NetworkX fails for rare pairs, the code falls back to a direct [u,v] edge; consider handling that carefully in future work.)

# 3) Experiments

## 3.1 Experiment 1 — Clean Single-Path (South → North)

**Intent.** Create a visually clean route that largely hugs the *shortest S→N street path* (the “spine”), visiting only regions that lie close to it.

### Steps.

1. Compute the shortest S→N path on streets (the spine).
2. Build the quad-tree over all deliveries.
3. **Filter** regions whose centroids lie within a corridor band (e.g., 3 km) of the spine.
4. **Order** the kept regions by their projection along the spine (south→north).
5. For each region in order: connect to its entry; do regional NN tour; stitch back to the sequence.
6. Connect last region to N.

### Outputs.

- Route PNG per n (visited/unvisited shown), quad-tree PNG, and a scalability panel plotting time/regions/coverage/efficiency.

**Notes.** No budget enforcement is applied in this clean-path variant; distance is observed/reported, not constrained.

## 3.2 Experiment 2 — Multi-Depot Comparison

**Intent.** Compare performance and coverage across several start→end depot pairs using the *same clean-path method*.

**Pairs.** Example: (S→E), (C→E), (W→S), (E→W), (N→S).

**Steps.** For each pair: load the same delivery set, run clean-path routing, measure computation time, memory (approx. peak), distance, coverage.

**Outputs.**

- **Heatmap/table** comparing distance, time, memory, and deliveries covered per depot pair.
- **Bar chart** of number of deliveries covered per pair.
- Region-distribution chart (nearest-depot counts) for intuition.

**Notes.** Also uses no budget; it's a head-to-head profile of geometry and workload differences by origin/destination.

## 3.3 Experiment 3 — Priority-Only Region Selection (No Corridor)

**Intent.** Make **priority** drive which regions are visited, not proximity to the S→N spine. Still travel S→N for readability.

**Selection.**

- Score each region by **priority weight**:  $50 \times \#(P1) + 30 \times \#(P2) + 20 \times \#(P3)$ .
- Filter out small regions: keep only those with  $\geq \text{min\_points\_per\_region}$ .
- Keep top-K regions by *(score, size)*; **no corridor distance** considered.
- Order selected regions south→north by spine projection; visit all points in each (regional NN) and stitch with street paths.

**Visuals.**

- **Priority-colored scatter** of all points (P1=red, P2=orange, P3=green).
- **Selected region overlays** (rectangles + centroids) on the same map.
- Standard route plot (visited vs unvisited) for each config.

**Knobs.** `min_points_per_region`, `top_k_regions`, `quad-tree max_points_per_region`.

## 4) Time & space complexity (practical big-O)

Let  $|V|$ =#nodes,  $|E|$ =#edges of the street graph;  $N$ =deliveries;  $R$ =regions;  $P_i$ =points in region  $i$  ( $\sum P_i = N$ ).

### Preprocessing

- KD-tree build for node snapping:  $O(|V| \log |V|)$  (once).
- Shortest S→N spine:  $O(|E| \log |V|)$  (once).
- Quad-tree:  $\sim O(N \log N)$ .

### Per-region entry selection

- Use **one** single-source Dijkstra from the current node:  $O(|E| \log |V|)$ ; then take  $\text{argmin}$  over region nodes (constant factor per region).

### Intra-region NN

- For  $P_i$  points:  $P_i$  iterations, each performs one Dijkstra from the current node  $\Rightarrow O(P_i \cdot |E| \log |V|)$  per region; summed:  $O(N \cdot |E| \log |V|)$ .

### Stitching between regions

- A handful of shortest paths between consecutive regions and to the end depot  $\Rightarrow O(R \cdot |E| \log |V|)$ .

### Total time

- Dominated by many SSSPs:  $\approx O(N \cdot |E| \log |V|)$ .

### Space

- Graph & KD-tree:  $O(|V| + |E|)$ .
- Deliveries/regions:  $O(N + R)$ .
- **Dijkstra cache** (LRU of single-source distance dicts): worst-case large; each cached entry stores distances to many nodes. Keep `maxsize` moderate and monitor resident memory.

**Empirical tip.** Runtime grows mainly with (*#regions visited*)  $\times$  (*avg region size*) and the density of the street graph. The single-Dijkstra entry selection keeps constants much lower than running Dijkstra per candidate.

## 5) Plot & file glossary (what each artifact means)

Artifact / Pattern	Produced by	What it shows	How to read	Notes/Caveats
Quad_Tree_Decomposition_Size_{n}.png	Exp 1	Quad-tree rectangles + all deliveries	Regions' spatial footprint and depth; centroids marked	Helps see clustering and region sizes
Clean_Path_Route_Size_{n}.png	Exp 1	Final S→N route on streets; visited vs unvisited points	Title includes total distance and coverage %	Corridor band implicit (not drawn)
Clean_Path_Scalability_Analysis.png	Exp 1	4 panels: time vs n, #regions vs n, coverage %, efficiency	Expect time and #regions to grow with n; coverage plateaus reveal corridor tightness	Efficiency = deliveries/km
experiment_2_results.csv	Exp 2	Numeric metrics per depot pair	Distance/time/memory/coverage per pair	Useful for tables in reports Values are normalized by colormap, exact text overlays included
Multi_Depot_Performance_Heatmap.png	Exp 2	Heatmap (distance, time, memory, covered)	Brighter/larger values per metric per depot pair	Sensitive to corridor width & city geometry Heuristic by haversine, not network distance
Delivery_Points_Covered_by_Depot_Pair.png	Exp 2	Bar chart of covered deliveries	Compare which OD pairs are most effective per pair	
Regional_Distribution_Analysis.png	Exp 2	Deliveries by nearest depot (bar chart)	Which depot "owns" more nearby points	
Priority_All_Points.png	Exp 3	All points colored by priority	Red=high, orange=med, green=low	No regions yet; useful sanity check

Artifact / Pattern	Produced by	What it shows	How to read	Notes/Caveats
Priority_Selected_Regions_{cfg}.png	Exp 3	Priority-colored points + <b>only selected</b> region rectangles	Confirms which regions the selector kept	Selection is by score & size only (no corridor)
Priority_NoCorridor_Route_{cfg}.png	Exp 3	Street route for priority-onl	Coverage tilts toward high-priority clusters	Distance may be larger than Exp 1
Priority_NoCorridor_Coverage_Comparison.png	Exp 3	Bar chart of total deliveries covered per config	Compare strict vs moderate selectors	Use with distance to see trade-offs
priority_no_corridor_results_sn.csv	Exp 3	Metrics per config (distance, coverage, efficiency)	Join with visuals for narrative	File naming indicates “no corridor” run

## 6) Interpretation guide

- **Coverage (%)**: visited deliveries  $\div$  total deliveries. Higher is better, but watch the corresponding **distance (km)**.
- **Efficiency (deliveries/km)**: how many deliveries you get per km of street travel. Useful to compare configurations independent of absolute scale.
- **#Regions used**: rough proxy for how fragmented the route becomes. Too many small regions can inflate stitch costs; too few big regions can cause detours.

## 7) Known limitations & future improvement ideas

- **Corridor (Exp 1/2)** may ignore cheap high-value clusters slightly off the spine. Option: allow soft scoring that trades detour for high priority.
- **Centroid distance** can misrepresent elongated regions. Option: min distance from **any** region point or convex hull to spine.
- **Nearest-neighbor tours** are fast but suboptimal. Option: a tiny **2-opt** pass per region (often saves 10–25% intra-region distance).

- **Feasibility/budget** is not enforced in these experiments; add an end-reachability reserve check if needed for production constraints.
  - **Dijkstra cache** memory: large LRU size can blow up resident memory on big graphs; trim the cache or cache only distances to delivery/depots.
  - **Fallback  $u \rightarrow v$  hop** on exceptions can draw a straight segment; add connectivity checks or retries on the giant component.
- 

## 8) Repro & tuning checklist

- CSV has `priority` with values in {1,2,3}. If missing, the loader defaults to 2 (neutral).
- For **Exp 1**: tune `corridor_width_km` (narrower = cleaner path, potentially less coverage).
- For **Exp 2**: same method as Exp 1; compare OD pairs fairly by keeping delivery set and parameters fixed.
- For **Exp 3**: set `min_points_per_region`, `top_k_regions` to actually bias toward high-priority clusters; no corridor is used.